

228.371 - Statistical Modelling for Engineers and Technologists

Week 8. Logistic Regression

Dr. Daniel Walsh
IIMS 3.07 x 41032
d.c.walsh@massey.ac.nz

Semester One - 2015

Logistic Regression

- ▶ Linear regression was concerned with modelling a *continuous* response variable with continuous or discrete predictor variables.
- ▶ Logistic regression is concerned with modelling a *dichotomous* (or binary) response variable with continuous or discrete predictor variables.

Example: fir.txt


Records details of 659 *balsam fir* trees from Boundary Water Canoe area wilderness, the tree's diameter, and whether or not the tree blow down in a 1999 storm with 90mph winds (1 indicates tree was blown down, 0 not).

```
fir <- read.table (file="Data/fir.txt", header=TRUE)
summary (fir)
```

dia		y	
Min.	: 5.0	Min.	:0.0000
1st Qu.:	6.0	1st Qu.:	0.0000
Median	: 8.0	Median	:0.0000
Mean	: 9.7	Mean	:0.3536
3rd Qu.:	12.0	3rd Qu.:	1.0000
Max.	:32.0	Max.	:1.0000


Example: fir.txt

- Proportion of tree blown down increases with diameter.

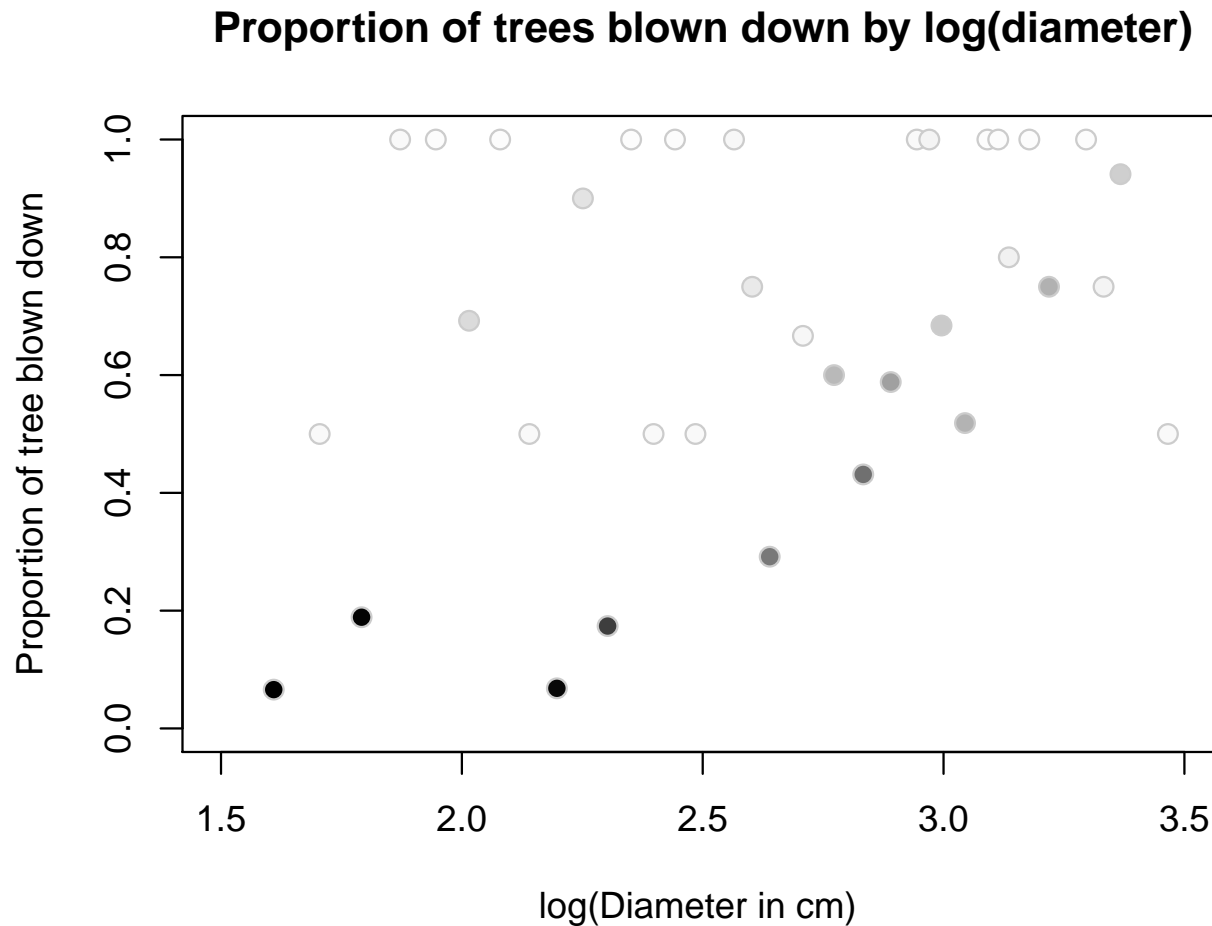


Example: fir.txt

- ▶ Diameter is left skewed - Log transform is taken.



Example: fir.txt



- ▶ Color indicates number of observations (the lighter the fewer).
- ▶ Moderate positive trend present.

Generalised linear model

The linear regression model assumes that the **mean** of a variable is linearly related to the predictors:

$$\mathbb{E}(y_i) = \beta_0 + \beta_1 x_{1i} + \cdots + \beta_{qi}$$

In a **generalised** linear model, a **function** of the mean is linearly related to the predictors:

$$g(\mathbb{E}(y_i)) = \beta_0 + \beta_1 x_{1i} + \cdots + \beta_{qi}$$

The function $g(\cdot)$ is called the **link** function.

Logistic regression


- ▶ When y is a dichotomous (0/1) variable, $\mathbb{E}(y_i)$ is the proportion (p) of 1's (successes). We want a function that will map p falling into the (0,1) interval (LHS of the equation) to the real line (RHS of the equation).

$$g(\mathbb{E}(y_i)) = \beta_0 + \beta_1 x_{1i} + \cdots + \beta_{qi}$$

- ▶ The most common choice of link, and the one we will focus on, is the *logit* function:

$$\text{logit}(p) = \log \left(\frac{p}{1-p} \right)$$

Logit function



$$\lim_{p \rightarrow 1} \text{logit}(p) = \infty$$

$$\lim_{p \rightarrow 0} \text{logit}(p) = -\infty$$

Logistic Regression in R

```
m1 <- glm (y ~ log(dia), data=fir, family=binomial)
summary (m1)
```

Call:

```
glm(formula = y ~ log(dia), family = binomial, data = fir)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.5073	-0.7565	-0.4936	0.8096	2.3272

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-7.8925	0.6325	-12.48	<2e-16 ***
log(dia)	3.2643	0.2761	11.82	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 856.21 on 658 degrees of freedom
Residual deviance: 655.24 on 657 degrees of freedom
AIC: 659.24

Number of Fisher Scoring iterations: 4

Where do the fitted β s come from?

- ▶ The β 's are chosen to maximise the probability of the observed data:

$$P(\mathbf{y} \mid \boldsymbol{\beta}) = \prod_{i=1}^n p_i^{y_i} (1 - p_i)^{1-y_i}$$

$$\text{logit}(p_i) = \beta_0 + \beta_1 x_{1i} + \cdots + \beta_q x_{qi}$$

- ▶ This is called “maximum likelihood.”
- ▶ We estimate the parameters ($\hat{\beta}_i$'s) so as to maximise the probability (or likelihood) of the observed data (\mathbf{y}).
- ▶ This is an iterative procedure.

Interpretation

Note:

$$\mu = \log \left(\frac{p}{1-p} \right) \Rightarrow p = \frac{e^\mu}{1 + e^\mu}$$

As in linear regression, non-zero β_j suggests x_j is a useful predictor. Suppose dia=5 hence $\log(\text{dia})=1.61$.

Then the probability the tree has blown down is:

```
exp(-7.8925+3.2643*log(5))/(1+exp(-7.8925+3.2643*log(5)))
```

```
[1] 0.06668194
```

Alternatively:

```
mu <- predict.glm (m1, newdata=data.frame(dia=5))
```

```
mu
```

```
[1] -2.638832
```

```
exp(mu)/(1+exp(mu))
```

```
[1] 0.06668071
```

Interpretation


However, the relationship between changes in x ($\log(\text{diameter})$) and changes in $\mathbb{E}(y)$ is more complicated. e.g. for $x = 2.5, 3.0, 3.5$

```
mu <- predict.glm (m1, newdata=data.frame(dia=exp(c(2.5,3.0,3.5))))  
exp(c(2.5,3.0,3.5))  
[1] 12.18249 20.08554 33.11545  
mu  
      1      2      3  
0.2681993 1.9003320 3.5324646  
exp(mu)/(1+exp(mu))  
      1      2      3  
0.5666508 0.8699291 0.9715975
```

Changing x by 0.5 (or any constant amount) has a different effect if we are changing it from 2.5 to 3.0 or 3.0 to 3.5.

Predicted probabilities

Probability tree is blown down



Odds and Log Odds

What is increasing linearly is the *log odds*: (Note: we have remove the observation subscript i for simplicity).

$$\log \left(\frac{p}{1-p} \right) = \beta_0 + \beta_1 x_1 + \cdots + \beta_q x_q$$

The odds,

$$\frac{p}{1-p},$$

are multiplied by a factor of $\exp(\beta_j)$ for each unit of increase of x_j .

Binned Data

We can also analyse the data by binning the response variable by the predictor.

```
successes <- table(fir$dia, fir$y)[,2:1]
uniqueD <- sort(unique(fir$dia))

cbind (successes, uniqueD)
      1  0 uniqueD
5      6 82      5.0
5.5    1  1      5.5
6      6 85      6.0
6.5    1  0      6.5
7     17 73      7.0
7.5    1  0      7.5
8     12 57      8.0
8.5    1  0      8.5
9     14 34      9.0
9.5    2  1      9.5
10     22 29     10.0
10.5    1  1     10.5
11     20 14     11.0
11.5    3  0     11.5
12     15 10     12.0
13     21  7     13.0
13.5    2  0     13.5
14     14 13     14.0
15     16  1     15.0
16     13  6     16.0
17      9  4     17.0
18      9  1     18.0
19      4  1     19.0
19.5    1  0     19.5
20      6  2     20.0

m2 <- glm(successes~log(uniqueD), family=binomial)
summary (m2)

Call:
glm(formula = successes ~ log(uniqueD), family = binomial)

Deviance Residuals:
      Min       1Q   Median       3Q      Max
-1.8980  -0.8097   0.3527   1.1355   2.3303

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept)   -7.8925     0.6325  -12.48  <2e-16 ***
log(uniqueD)    3.2643     0.2761   11.82  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

      Null deviance: 250.856  on 34  degrees of freedom
Residual deviance:  49.891  on 33  degrees of freedom
AIC: 117.52

Number of Fisher Scoring iterations: 4
```


What has and hasn't changed?

Unchanged:

- ▶ Estimates and Standard errors
- ▶ Difference between null and residual Deviance and corresponding degrees of freedom

Changed:

- ▶ Raw deviances/degrees of freedom
- ▶ Residuals

Deviance

Deviance is a measure of fit of the model to the data. Lower deviance means better fit. We can calculate by summing over observations (where y_i is 0/1 (success/failure) of observation i):

$$Deviance = \sum_i^n -2 \times [y_i \log(\hat{\mu}_i) + (1 - y_i) \log(1 - \hat{\mu}_i)]$$

Or by aggregate over observations with the same set of x -values (y_j is the number of successes and n_j is the number of observations, corresponding to the j th value of x) :

$$Deviance = \sum_1^m -2 \times [y_j (\log(\hat{\mu}_j) - y_j/n_j) + (n_j - y_j) (\log(1 - \hat{\mu}_j) - \log((n_j - y_j)/n_j))]$$

The null deviance (deviance under the null hypothesis) uses a single μ ; the residual deviance uses the μ given by our model.

Model Utility

We want to test whether our predictions are better with predictor x .

- ▶ Null hypothesis: Null model (each observation has same probability of success) is adequate;
- ▶ Test statistic: difference in deviances.
- ▶ Null distribution of the test statistic: Chi-squared distribution with degrees of freedom, $df = df_{Null} - df_{model}$.

For the `fir.txt` data:

Deviances are given when model is printed. The p -value to reject null model can be calculated with `pchisq()`.

`m2`

```
Call: glm(formula = successes ~ log(uniqueD), family = binomial)
```

```
Coefficients:
```

```
(Intercept)  log(uniqueD)
      -7.892         3.264
```

```
Degrees of Freedom: 34 Total (i.e. Null);  33 Residual
```

```
Null Deviance:      250.9
```

```
Residual Deviance: 49.89  AIC: 117.5
```

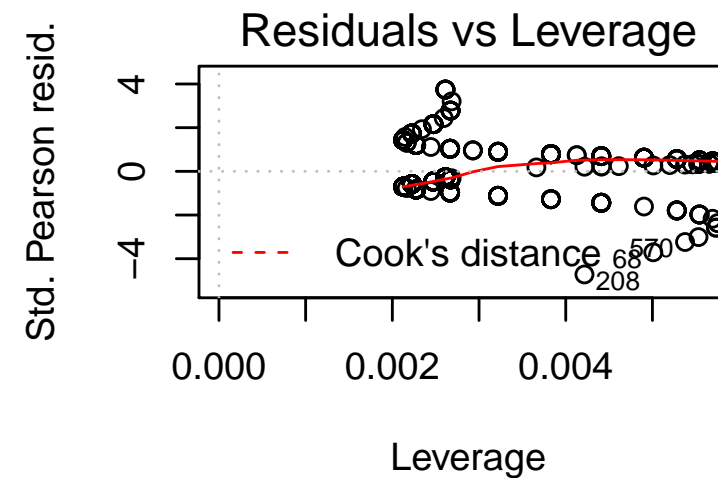



```
## P-value
```

```
pchisq(q=250.856-49.891, df=34-33, lower=FALSE)
```





```
[1] 1.286026e-45
```

- ▶ This is the same regardless of how we fit the model (`m1` vs `m2`).
- ▶ Unlike linear regression, it is not identical to the z -test shown for the coefficient of x .

Diagnostic Plots - Individual Observations



Diagnostic Plots - Aggregated Data



Residuals

- Deviance residuals:

$$\sqrt{-2 \left[y_j \left(\log(\hat{\mu}_j) - \frac{y_j}{n_j} \right) + (n_j - y_j) \left(\log(1 - \hat{\mu}_j) - \log\left(\frac{n_j - y_j}{n_j}\right) \right) \right]}$$

with the usual sign.

- Pearson Residuals:

$$r_p = \frac{y_i - \hat{\mu}_i}{\sqrt{\mu_i(1 - \mu_i)/n_i}}$$

- Fitted values plotted are $\text{logit}(\hat{\mu}_i) = \beta_o + \beta_x x_{1i} \dots$ (logit scale not probability scale)

We are looking for:

- ▶ No outliers.
- ▶ No trends.
- ▶ Standardised Pearson's residual may have skewed distribution, but should be free of trends.
- ▶ No points of high influence (outside Cook's distance contours).
- ▶ Cook's distance and leverage interpreted the same way as in linear regression; calculated differently (accounting for differing variances). No longer extreme x-values that are most influential.

Model Adequacy

- ▶ 'Saturated model' uses the sample proportion for each unique set of x values. Deviance, degrees of freedom are zero.
- ▶ Only sensible when there are multiple observation per predictor set.
- ▶ Can use deviance to see if this improves on fitted model.
- ▶ If not, our model fits as well as can be hoped.
- ▶ E.g. for the tree fall data, residual deviance for logistic regression with a single predictor is 49.891.

```
pchisq(49.891-0, 33-0, lower=FALSE)
```

```
[1] 0.0299041
```

We do not have a “well fitting model” just yet.

Multiple Logistic Regression

- ▶ The various complications that we considered for linear regression are also possible for logistic regression.
- ▶ Consider an enlarged tree data set:
 - ▶ More trees-two species, Balsam Fir and Blue Spruce
 - ▶ Measurement of “local wind severity” - proportion of trees in that area knocked down.
 - ▶ Consider interactions/multiplicative relationships

Example: firspruce.txt

The larger data set contains two species of tree.

```
firspruce <- read.table (file="Data/firspruce.txt", header=TRUE)
summary (firspruce)
```

dia	severity	y	species
Min. : 5.00	Min. :0.02421	Min. :0.0000	fir :659
1st Qu.: 8.00	1st Qu.:0.21894	1st Qu.:0.0000	spruce:970
Median :11.50	Median :0.40087	Median :0.0000	
Mean :12.31	Mean :0.41698	Mean :0.4696	
3rd Qu.:15.00	3rd Qu.:0.59286	3rd Qu.:1.0000	
Max. :51.00	Max. :0.98327	Max. :1.0000	

Example: firspruce.txt: Backwards Stepwise Selection

```
m3 <- step(glm(y ~ log(dia) * species * severity,  
              family='binomial', data=firspruce))
```

Start: AIC=1485.73

```
y ~ log(dia) * species * severity
```

	Df	Deviance	AIC
- log(dia):species:severity	1	1470.8	1484.8
<none>		1469.7	1485.7

Step: AIC=1484.8

```
y ~ log(dia) + species + severity + log(dia):species + log(dia):severity +  
    species:severity
```

	Df	Deviance	AIC
- species:severity	1	1470.8	1482.8
<none>		1470.8	1484.8
- log(dia):species	1	1478.6	1490.6
- log(dia):severity	1	1512.3	1524.3

Step: AIC=1482.84

```
y ~ log(dia) + species + severity + log(dia):species + log(dia):severity
```

	Df	Deviance	AIC
<none>		1470.8	1482.8
- log(dia):species	1	1478.7	1488.7
- log(dia):severity	1	1517.0	1527.0

How do we picture what this model looks like?

Example: firspruce.txt: Model summary

```
summary (m3)
```

```
Call:
```

```
glm(formula = y ~ log(dia) + species + severity + log(dia):species +  
     log(dia):severity, family = "binomial", data = firspruce)
```

```
Deviance Residuals:
```

Min	1Q	Median	3Q	Max
-3.1952	-0.6564	-0.3659	0.6912	2.4035

```
Coefficients:
```

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-4.6969	1.0318	-4.552	5.31e-06	***
log(dia)	1.0537	0.4385	2.403	0.01627	*
speciesspruce	2.3090	0.8907	2.592	0.00953	**
severity	-8.1858	2.0918	-3.913	9.10e-05	***
log(dia):speciesspruce	-1.0459	0.3778	-2.768	0.00564	**
log(dia):severity	5.6687	0.8759	6.472	9.67e-11	***

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
(Dispersion parameter for binomial family taken to be 1)
```

```
Null deviance: 2252.3  on 1628  degrees of freedom  
Residual deviance: 1470.8  on 1623  degrees of freedom  
AIC: 1482.8
```

```
Number of Fisher Scoring iterations: 5
```

Prediction

Plot the relationship between probability of blow down and diameter; use colours/line types to show effect of Species and local severity.

```
DatFir <- data.frame (dia=rep(seq(5, 50, 1),2),
                      severity=rep(c(0.25, 0.75), each=46),
                      species=rep("fir", 46*2))
DatSpruce <-data.frame(dia=rep(seq(5, 50, 1),2),
                      severity=rep(c(0.25, 0.75), each=46),
                      species=rep("spruce", 46*2))

summary (DatFir)

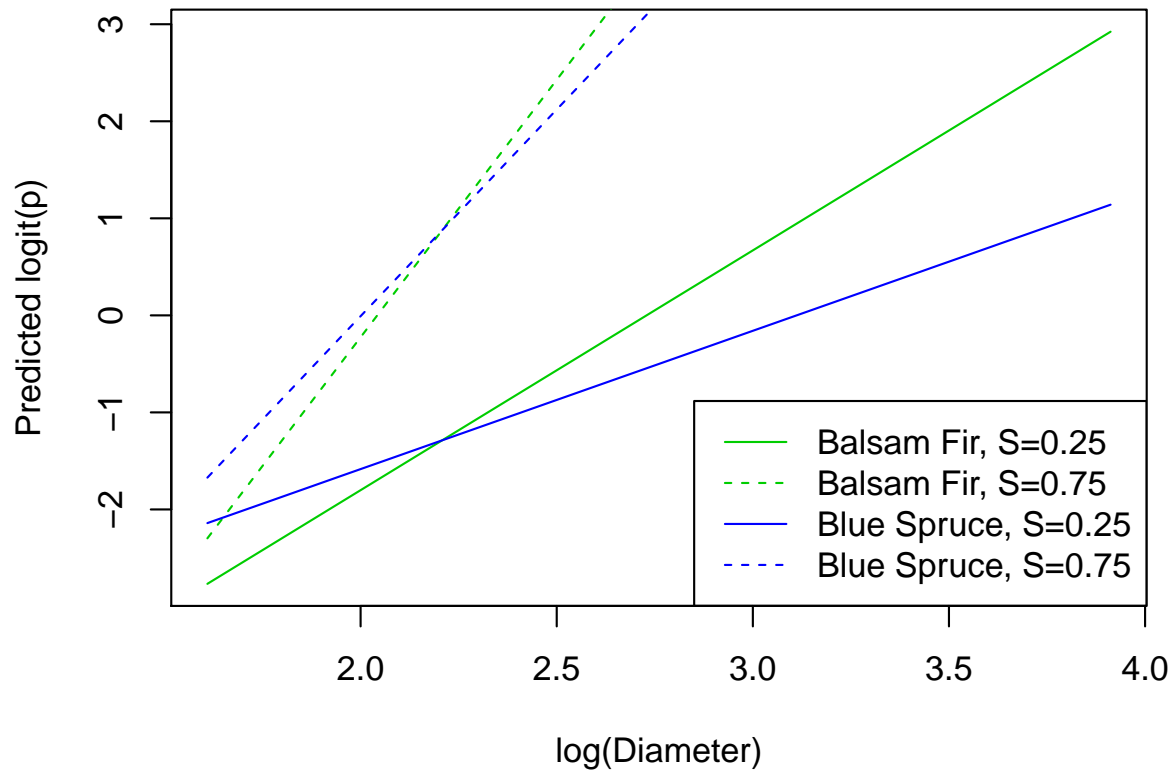
      dia      severity      species
Min.   : 5.0   Min.   :0.25   fir:92
1st Qu.:16.0   1st Qu.:0.25
Median :27.5   Median :0.50
Mean   :27.5   Mean   :0.50
3rd Qu.:39.0   3rd Qu.:0.75
Max.   :50.0   Max.   :0.75

summary (DatSpruce)

      dia      severity      species
Min.   : 5.0   Min.   :0.25   spruce:92
1st Qu.:16.0   1st Qu.:0.25
Median :27.5   Median :0.50
Mean   :27.5   Mean   :0.50
3rd Qu.:39.0   3rd Qu.:0.75
Max.   :50.0   Max.   :0.75


pred.Fir.lin      <- predict.glm (m3, DatFir)
pred.Fir.prob     <- predict.glm (m3, DatFir, type='response')
pred.Spruce.lin   <- predict.glm (m3, DatSpruce)
pred.Spruce.prob  <- predict.glm (m3, DatSpruce, type='response')
```

Plots produced




```
plot(log(DatFir$dia[1:46]),pred.Fir.lin[1:46], type="l",  
      xlab='log(Diameter)', ylab='Predicted logit(p)',col=3)  
lines(log(DatFir$dia[47:92]), pred.Fir.lin[47:92], lty=2,col=3)  
lines(log(DatSpruce$dia[47:92]),pred.Spruce.lin[47:92], lty=2,col=4)  
lines(log(DatSpruce$dia[1:46]), pred.Spruce.lin[1:46], col=4)  
legend("bottomright", lty=c(1,2,1,2), col=c(3,3,4,4),  
      legend=c("Balsam Fir, S=0.25", "Balsam Fir, S=0.75",  
                "Blue Spruce, S=0.25", "Blue Spruce, S=0.75"))
```

Plots produced




```
plot(log(DatFir$dia[1:46]),pred.Fir.prob[1:46], type="l",  
      ylim=c(0,1), xlab='log(Diameter)', ylab='Predicted Probability',col=3)  
lines(log(DatFir$dia[47:92]), pred.Fir.prob[47:92], lty=2,col=3)  
lines(log(DatSpruce$dia[47:92]),pred.Spruce.prob[47:92], lty=2,col=4)  
lines(log(DatSpruce$dia[1:46]), pred.Spruce.prob[1:46], col=4)  
legend("bottomright", lty=c(1,2,1,2), col=c(3,3,4,4),  
      legend=c("Balsam Fir, S=0.25", "Balsam Fir, S=0.75",  
                "Blue Spruce, S=0.25", "Blue Spruce, S=0.75"))
```


Ingots ready for rolling



Ingots ready for rolling



Example: `ingots.txt`

- ▶ Three different soaking times and five different heating times
- ▶ Between 7 and 56 observations per combination of factors
- ▶ Small number of “successes” (ingot not ready).

Example: `ingots.txt` - ANOVA

- ▶ Treat levels as factors to start with
- ▶ As with regular ANOVA we first want to assess whether there is an interaction-will use difference in deviance test
- ▶ Model with interaction always provides a perfect fit (saturated model)
- ▶ Between 7 and 56 observations per combination of factors
- ▶ Small number of “successes” (ingot not ready).

Example: ingots.txt - Data

```
ingots <- read.table (file="Data/ingots.txt", header=TRUE)
```

```
ingots
```

	HeatingTime	SoakingTime	NotReady	Ready
1	7	1.0	0	10
2	7	1.7	0	17
3	7	2.2	0	7
4	7	2.8	0	12
5	7	4.0	0	9
6	14	1.0	0	31
7	14	1.7	0	43
8	14	2.2	2	31
9	14	2.8	0	31
10	14	4.0	0	19
11	27	1.0	1	55
12	27	1.7	4	40
13	27	2.2	0	21
14	27	2.8	1	21
15	27	4.0	1	15
16	51	1.0	3	10
17	51	1.7	0	1
18	51	2.2	0	1
19	51	4.0	0	1

Example: ingots.txt - Anova

```
m4a <- glm (cbind(NotReady, Ready) ~ factor(SoakingTime) + factor(HeatingTime),  
            data=ingots, family="binomial")  
m4b <- glm (cbind(NotReady, Ready) ~ factor(SoakingTime) * factor(HeatingTime),  
            data=ingots, family="binomial")  
anova (m4a, m4b, test ="Chisq")
```

Analysis of Deviance Table

```
Model 1: cbind(NotReady, Ready) ~ factor(SoakingTime) + factor(HeatingTime)  
Model 2: cbind(NotReady, Ready) ~ factor(SoakingTime) * factor(HeatingTime)  
  Resid. Df Resid. Dev Df Deviance Pr(>Chi)  
1         11      11.274  
2          0         0.000 11   11.274  0.4206
```

Fit of additive model is OK.

Example: ingots.txt - Anova

```
anova (m4a, test="Chisq")
```

Analysis of Deviance Table

Model: binomial, link: logit

Response: cbind(NotReady, Ready)

Terms added sequentially (first to last)

	Df	Deviance	Resid.	Df	Resid.	Dev	Pr(>Chi)
NULL				18		25.395	
factor(SoakingTime)	4	1.0402		14		24.355	0.903633
factor(HeatingTime)	3	13.0810		11		11.274	0.004465 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

P-values will depend on order.

Example: ingots.txt - Anova

```
m4c <- glm( cbind(NotReady, Ready) ~  
            as.factor(HeatingTime)+as.factor(SoakingTime),  
            data=ingots, family="binomial")  
anova(m4c,test="Chisq")
```

Analysis of Deviance Table

Model: binomial, link: logit

Response: cbind(NotReady, Ready)

Terms added sequentially (first to last)

	Df	Deviance	Resid.	Df	Resid.	Dev	Pr(>Chi)				
NULL				18		25.395					
as.factor(HeatingTime)	3	12.7100		15		12.685	0.005308 **				
as.factor(SoakingTime)	4	1.4113		11		11.274	0.842234				

Signif. codes:	0	'***'	0.001	'**'	0.01	'*'	0.05	'.'	0.1	' '	1

P-values will depend on order.

Example: ingots.txt - Heat as a linear predictor

Heating Time is a (borderline) significant predictor; could be used as a linear predictor to gain a degree of freedom.

```
m5a <- glm(cbind(NotReady, Ready) ~ factor(HeatingTime), data=ingots, family="binomial")
m5b <- glm(cbind(NotReady, Ready) ~ HeatingTime, data=ingots, family="binomial")
anova(m5a, m5b, test="Chisq") ## Shows linear predictor works just as well
```

Analysis of Deviance Table

Model 1: cbind(NotReady, Ready) ~ factor(HeatingTime)

Model 2: cbind(NotReady, Ready) ~ HeatingTime

	Resid. Df	Resid. Dev	Df	Deviance	Pr(>Chi)
1	15	12.685			
2	17	13.782	-2	-1.0961	0.5781

```
anova(m5b, test="Chisq")      ## Shows p-value associated with linear predictor
```

Analysis of Deviance Table

Model: binomial, link: logit

Response: cbind(NotReady, Ready)

Terms added sequentially (first to last)

		Df	Deviance	Resid. Df	Resid. Dev	Pr(>Chi)
NULL				18	25.395	
HeatingTime	1	11.614		17	13.782	0.0006546 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Example: ingots.txt - Coefficients

We should also have a look at the actual coefficients to see if Heating Time is good or bad:

```
summary(m5b)

Call:
glm(formula = cbind(NotReady, Ready) ~ HeatingTime, family = "binomial",
    data = ingots)

Deviance Residuals:
    Min       1Q   Median       3Q      Max 
-1.27249  -0.79572  -0.51483  -0.09311   1.71887 

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept) -5.41518     0.72754  -7.443 9.83e-14 ***
HeatingTime  0.08070     0.02236   3.610 0.000307 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 25.395  on 18  degrees of freedom
Residual deviance: 13.782  on 17  degrees of freedom
AIC: 32.109

Number of Fisher Scoring iterations: 5
```

Longer Heating Time seems to increase the probability of the Ingot being not ready.

Another ANOVA example

Time and shoot length used in grafting plum trees:

length	time	success	trials
long	autumn	156	240
long	spring	84	240
short	autumn	107	240
short	spring	31	240

```
plum <- c(156,84, 107,31)
plum <- cbind(plum, 240-plum)
shoot <- factor(c("long", "long", "short", "short"))
time <- factor(c("fall", "spring", "fall", "spring"))
```

Another ANOVA example

```
m6 <- glm(plum ~ shoot*time, family="binomial")
anova(m6, test="Chisq")
```

Analysis of Deviance Table

Model: binomial, link: logit

Response: plum

Terms added sequentially (first to last)


	Df	Deviance	Resid.	Df	Resid. Dev	Pr(>Chi)					
NULL				3	151.019						
shoot	1	45.837		2	105.182	1.285e-11 ***					
time	1	102.889		1	2.294	< 2.2e-16 ***					
shoot:time	1	2.294		0	0.000	0.1299					

Signif. codes:	0	'***'	0.001	'**'	0.01	'*'	0.05	'.'	0.1	' '	1

No interaction term needed.

Interaction plot

If there are no groups with all successes or all failures, it can be useful to make an interaction plot on the logit scale.



```
logit <- function (p) { log(p/(1-p)) }  
interaction.plot (time, shoot, logit(plum[,1]/240),  
                  ylab="logit(probability of success)")
```

Which are the favourable conditions?

```
m7 <- glm (plum ~ time + shoot, family="binomial")
summary (m7)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	0.7138	0.1217	5.867	4.45e-09	***
timespring	-1.4275	0.1465	-9.747	< 2e-16	***
shootshort	-1.0177	0.1455	-6.995	2.64e-12	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Spring and short shoots decrease the probability of success.

Case Control

- ▶ Rather than sampling from the population, enrich the number of 'Cases' (usually units with problems)
- ▶ Useful when:
 - ▶ Cases are rare AND
 - ▶ Measuring the predictor variables is expensive (time, money, destructive testing).
- ▶ β_1 to β_k are the same (and more accurately estimated) as in the case where a population sample is taken. Useful for understanding what predictors are important
- ▶ β_0 is not the same... cannot compute probabilities that pertain to samples from the population.

An industrial example

- ▶ Taken from “Effective Monitoring of Processes with Parts Per Million Defective: A Hard Problem!”
- ▶ Exhaust valve seats force fitted by insertion into the head of an engine. Four per engine.
- ▶ Expected defective rate is less than 50 PPM (cases rare).
- ▶ Quality of fit judged by visual inspection using feeler gauges (inspection expensive).
- ▶ 25 defective seat insertions collected over time; other three in each engine used as controls. (“Pareto analysis” showed no evidence poor fit depends on location.) 100 observations total.

An industrial example

- ▶ Eleven explanatory variables (Measurements of force taken during the automated insertion process, dimensional and physical characteristics of the valve seat and pocket in the head where the seat is inserted).
- ▶ Logistic regression used to predict defectives; three variables identified as important predictors (Inference focused on $\beta_1 \dots$)
- ▶ Quality control began charting these three continuous variables (x-bar chart, s-chart) to improve the quality of this part of the process.

Example with data

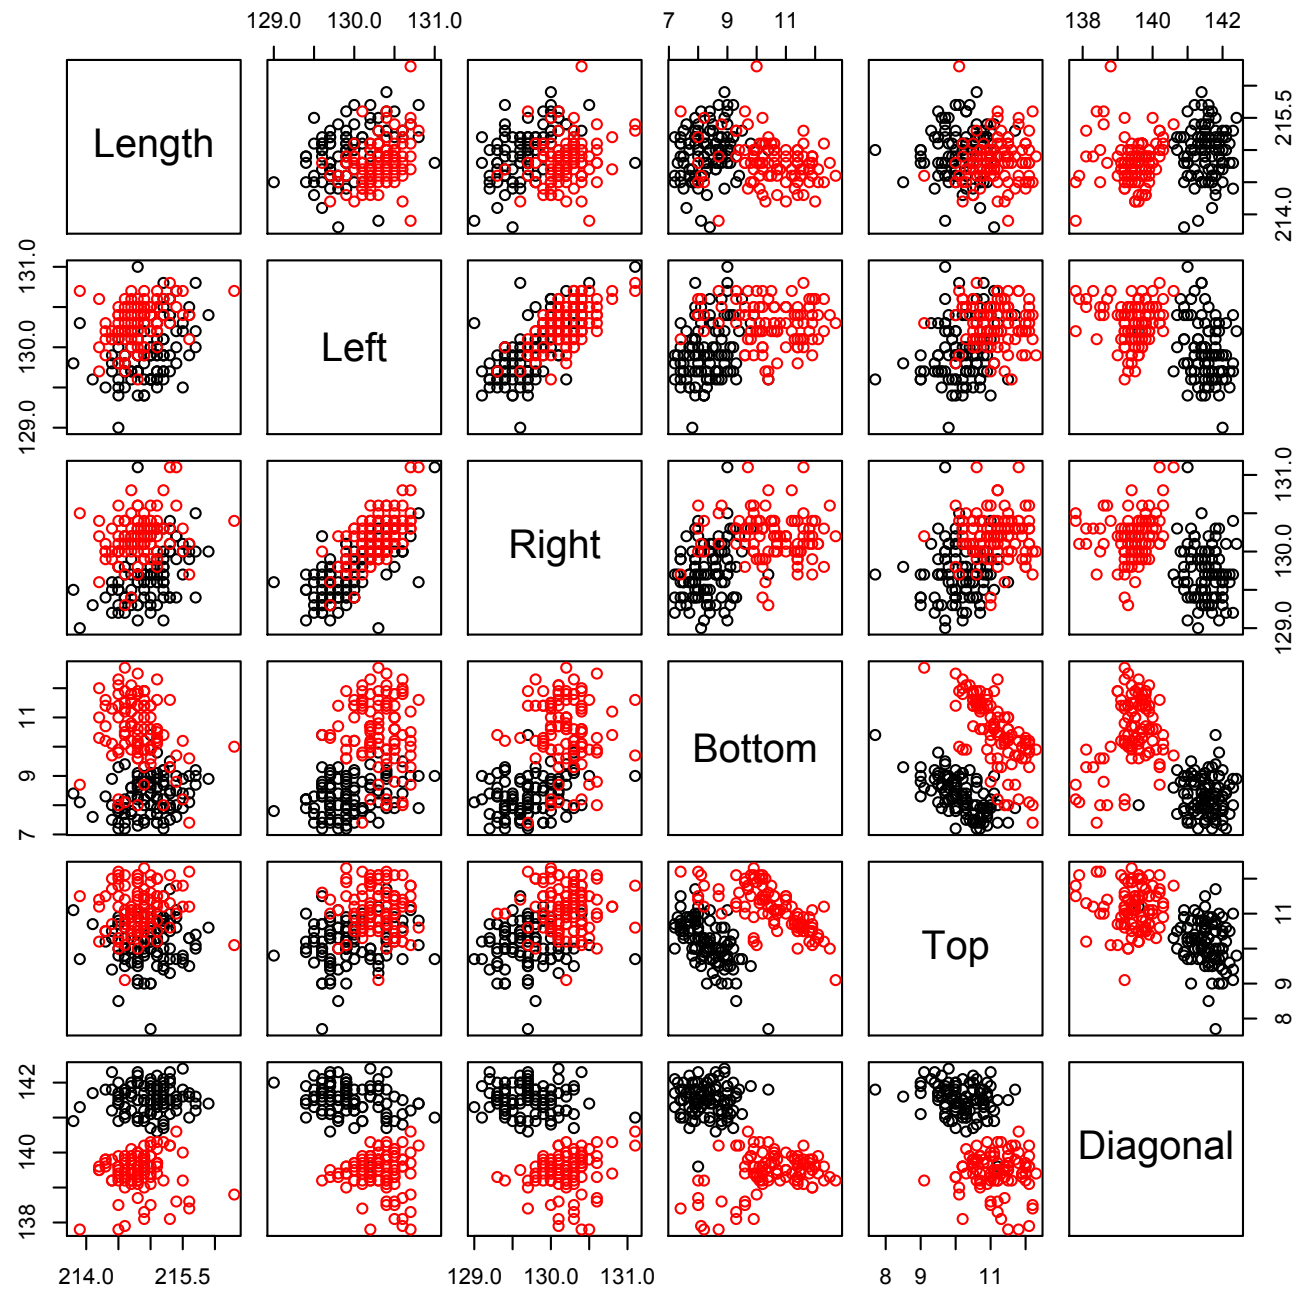
- ▶ 100 counterfeit and 100 legitimate swiss banknotes.
- ▶ 6 measurements (Length, Width, Diagonal, border widths) to use as predictors

```
banknote <- read.table ("Data/banknote.txt", header=TRUE)  
summary (banknote)
```

Length	Left	Right	Bottom
Min. :213.8	Min. :129.0	Min. :129.0	Min. : 7.200
1st Qu.:214.6	1st Qu.:129.9	1st Qu.:129.7	1st Qu.: 8.200
Median :214.9	Median :130.2	Median :130.0	Median : 9.100
Mean :214.9	Mean :130.1	Mean :130.0	Mean : 9.418
3rd Qu.:215.1	3rd Qu.:130.4	3rd Qu.:130.2	3rd Qu.:10.600
Max. :216.3	Max. :131.0	Max. :131.1	Max. :12.700

Top	Diagonal	Y
Min. : 7.70	Min. :137.8	Min. :0.0
1st Qu.:10.10	1st Qu.:139.5	1st Qu.:0.0
Median :10.60	Median :140.4	Median :0.5
Mean :10.65	Mean :140.5	Mean :0.5
3rd Qu.:11.20	3rd Qu.:141.5	3rd Qu.:1.0
Max. :12.30	Max. :142.4	Max. :1.0

Predictor Variables



Model Fitting

- ▶ Model fitting algorithm does not converge for the model with all 6 predictors.
- ▶ Try forward stepwise. Model with Diagonal and Bottom selected.

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	99422.9	5433597.5	0.018	0.985
Diagonal	-751.8	41093.2	-0.018	0.985
Bottom	688.7	37796.3	0.018	0.985

- ▶ Perfect separation, unstable fit.

Model Fitting

- ▶ Try the two variables that give the best visual separation in the pairs plot, Top and Diagonal.

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	884.197	215.637	4.100	4.12e-05	***
Diagonal	-6.450	1.544	-4.178	2.94e-05	***
Top	2.016	1.326	1.520	0.129	

- ▶ Not clear 'top' is helping. Model with just Diagonal:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	989.547	219.465	4.509	6.52e-06	***
Diagonal	-7.044	1.562	-4.508	6.54e-06	***

Model Fitting

- ▶ How to visualise the difference?.
- ▶ $p=0.5$ would be a reasonable place to change decision about the note being real/counterfeit. What does this look like in predictor space?
- ▶ $p=0.5$ means $\text{logit}(p)=0$.

$$0 = 884.197 - 6.450 \times \textit{Diagonal} + 2.016 \times \textit{Top};$$

rearranging we get


$$\textit{Top} = -438.5898 + 3.199405 \times \textit{Diagonal}.$$

- ▶ For the single variable model:

$$0 = 989.547 - 7.044 \times \textit{Diagonal}$$

giving $\textit{Diagonal} = 140.48$. Warnings about “fitted probabilities of 0 or 1.”

Model Fitting



Back in real life...

When counterfeit bills constitute 1% of all bills

